

Scalable Social Feed System Design

Executive Summary

This document presents a comprehensive architectural solution for building a social media feed system capable of handling millions of users and posts. The design addresses the critical "celebrity problem" where users with massive follower counts can overwhelm traditional systems. Through an innovative hybrid approach combining push and pull strategies, the system maintains consistent sub-second response times regardless of user scale.

Problem Statement

Social media platforms face significant scalability challenges as they grow. The most critical challenge occurs when users with millions of followers create posts, triggering what is known as the "celebrity problem." In traditional fan-out-on-write systems, a single post from a celebrity user can cause:

- **Massive Write Amplification:** Millions of database writes occurring simultaneously
- **System Overload:** Message queues and databases becoming overwhelmed
- **Degraded Performance:** Increased latency affecting all platform users
- **Resource Exhaustion:** Computational resources consumed by a single operation

These issues not only affect the celebrity users but degrade the experience for the entire user base, making a scalable solution essential.

Proposed Solution

The solution employs a hybrid architecture that intelligently routes content based on user characteristics. This approach combines the best aspects of both push and pull models:

Dual Strategy Approach

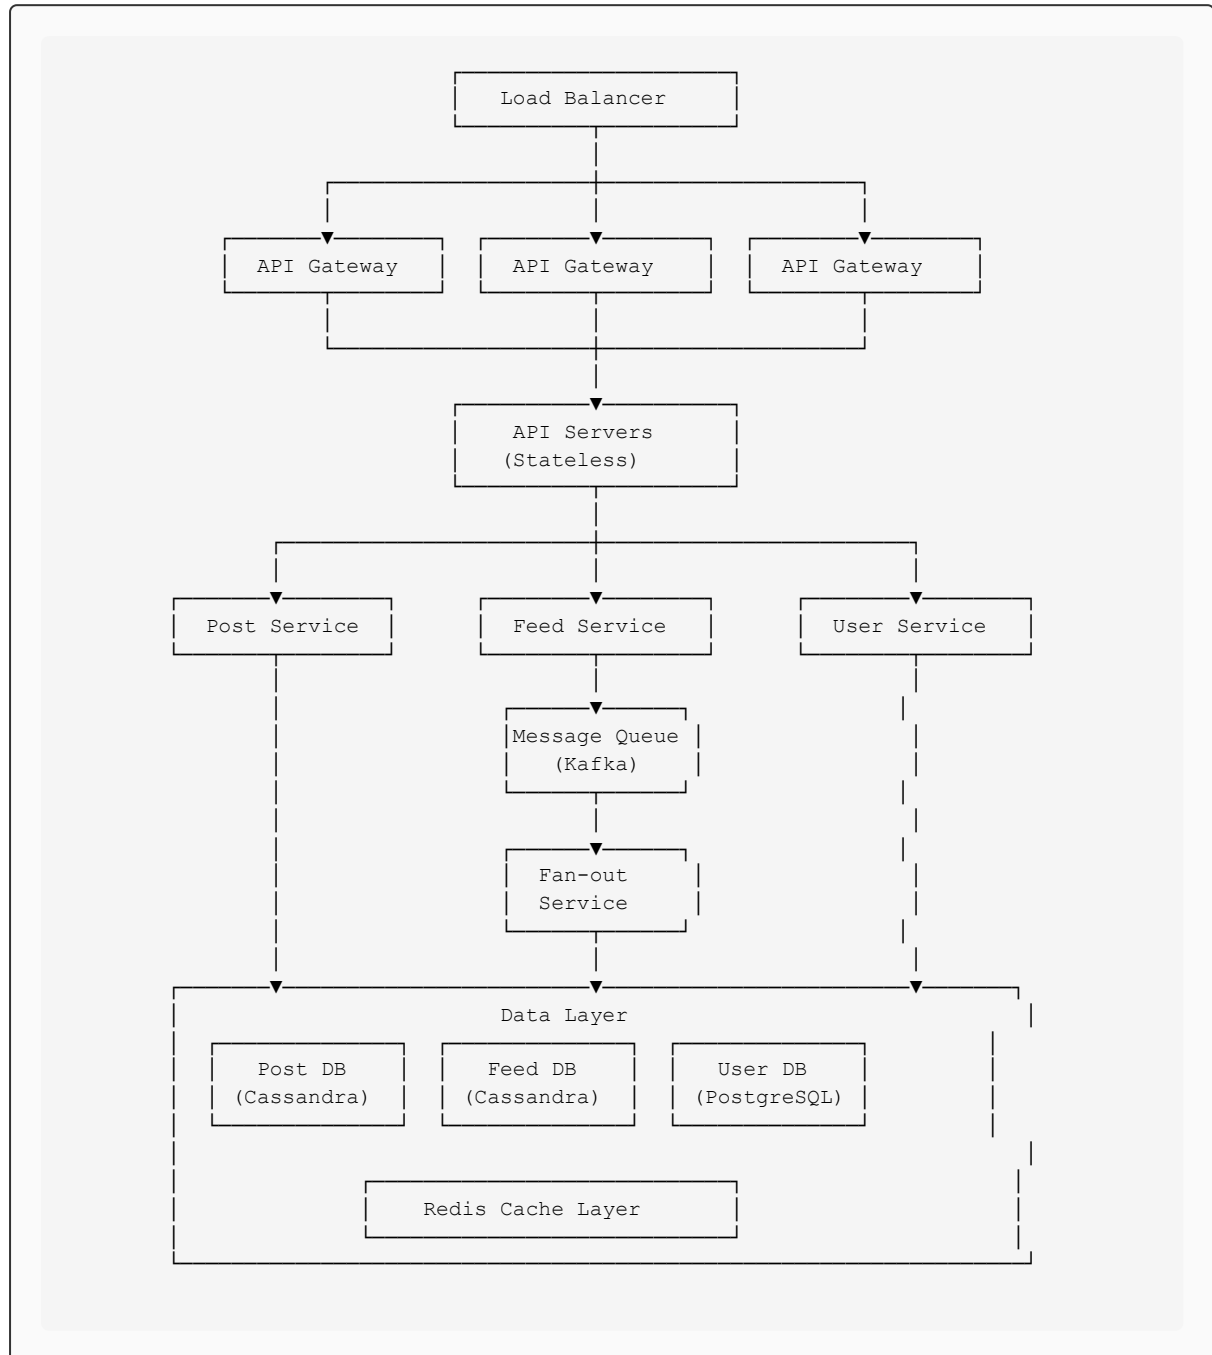
- **Push Model (Fan-out on Write):** Applied to regular users with fewer than 5,000 followers. Posts are immediately propagated to follower feeds, ensuring low-latency access.
- **Pull Model (Fan-out on Read):** Applied to celebrity users with 5,000 or more followers. Posts are stored once and retrieved on-demand during feed generation.

This threshold-based routing ensures optimal performance across all user segments while preventing system overload.

System Architecture

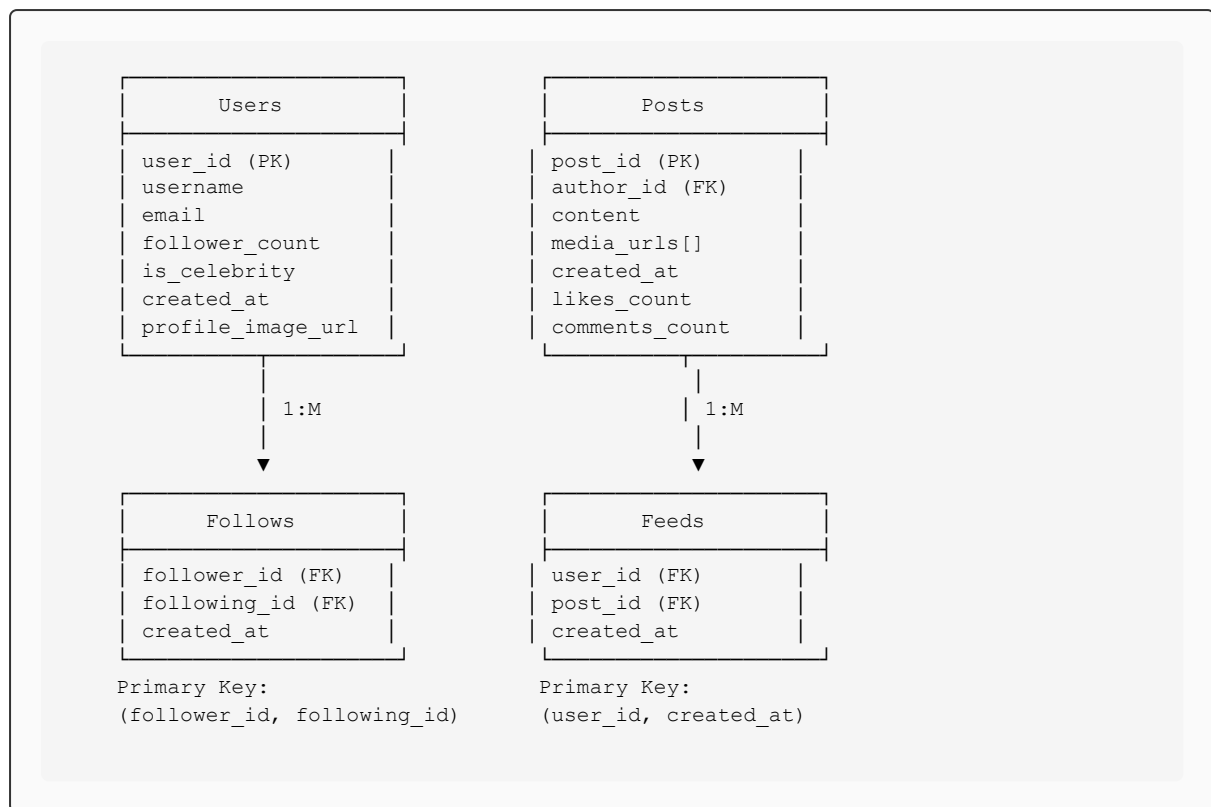
High-Level System Design

The architecture follows a microservices pattern with clear separation of concerns:



Data Model Design

The data model is optimized for both write and read operations:



Core Components

1. API Gateway Layer

- Handles request routing and load distribution
- Implements authentication and authorization
- Enforces rate limiting to prevent abuse
- Provides request/response transformation

2. Application Services

- **Post Service:** Manages post creation, storage, and retrieval
- **Feed Service:** Assembles and delivers personalized feeds
- **User Service:** Handles user profiles and relationship management
- **Fan-out Service:** Processes asynchronous feed updates for regular users

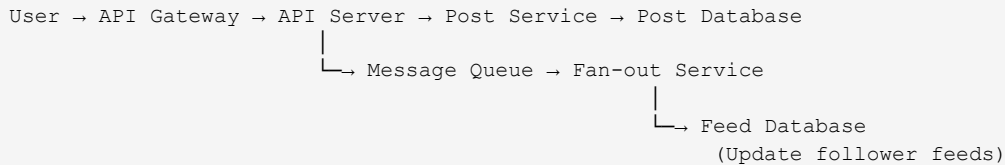
3. Data Storage Layer

- **PostgreSQL:** Stores user data requiring ACID properties
- **Cassandra:** Handles high-volume post and feed data
- **Redis:** Provides caching for frequently accessed content
- **Kafka:** Manages asynchronous message processing

Data Flow Analysis

Write Path - Regular User Flow

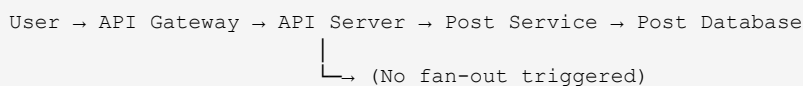
When a regular user creates a post, the system follows this sequence:



1. User submits post through client application
2. API Gateway validates and forwards request
3. API Server checks user type (`follower_count < 5,000`)
4. Post is saved to Post Database
5. Message is published to Kafka queue
6. Fan-out Service processes message asynchronously
7. Follower feeds are updated in batches

Write Path - Celebrity User Flow


Celebrity posts follow a simplified path:



1. Celebrity user submits post
2. API Gateway validates and forwards request
3. API Server detects celebrity status (`follower_count ≥ 5,000`)
4. Post is saved to Post Database only
5. No fan-out process is initiated

Read Path - Hybrid Feed Assembly

Feed generation combines pre-computed and on-demand content:



1. User requests their feed
2. API Server initiates parallel queries:
 - Fetch pre-computed feed from Cache/Feed Database
 - Query recent posts from followed celebrities
3. Results are merged and sorted by timestamp
4. Final feed is returned to user

Performance Analysis

Write Performance Metrics

- **Regular Users:** $O(1)$ database write + $O(N)$ asynchronous fan-out
- **Celebrity Users:** $O(1)$ database write only
- **Write Latency:** $< 50\text{ms}$ for all user types
- **Throughput:** 1M+ posts per minute

Read Performance Metrics

- **Cache Hit Rate:** $> 95\%$ for active users
- **Feed Generation:** $< 100\text{ms}$ (99th percentile)
- **Celebrity Post Query:** $O(M)$ where M = followed celebrities
- **Concurrent Users:** 10M+ simultaneous connections

Scalability Characteristics

- Horizontal scaling through service replication
- Database sharding based on `user_id`
- Geographic distribution via CDN integration
- Elastic resource allocation during peak loads

Implementation Guidelines

Technology Stack Recommendations

- **Backend Services:** Java/Spring Boot or Go for high performance
- **API Gateway:** Kong or AWS API Gateway
- **Message Queue:** Apache Kafka for reliability and scale
- **Caching:** Redis Cluster with persistence
- **Monitoring:** Prometheus + Grafana for metrics
- **Tracing:** Jaeger for distributed tracing

Optimization Strategies

1. **Celebrity Post Caching:** Cache celebrity posts with longer TTL values
2. **Batch Processing:** Group fan-out operations to reduce database load
3. **Connection Pooling:** Optimize database connections across services
4. **Compression:** Use protocol buffers for inter-service communication
5. **Adaptive Thresholds:** Dynamically adjust celebrity threshold based on system load

Monitoring and Operations

- Real-time dashboards for system health
- Automated alerts for performance degradation
- Capacity planning based on growth projections
- A/B testing framework for threshold optimization
- Disaster recovery with multi-region failover

Conclusion

This hybrid architecture provides a robust solution to the scalability challenges faced by modern social media platforms. By intelligently routing content based on user characteristics, the system achieves:

- Consistent performance regardless of user follower counts
- Efficient resource utilization through selective fan-out
- Seamless scaling to support hundreds of millions of users
- Maintainable codebase through clear separation of concerns

The design ensures that platform growth does not compromise user experience, making it suitable for both startup and enterprise-scale deployments.